



1 Movie Recommendation System

- Student name: Vi Bui
- Student pace: Part-Time
- Scheduled project review date/time: Thur. 03/03/22 (Non-Technical); Tues. 03/08/22 (Technical)
- Instructor name: Claude Fried
- Blog post URL: <https://datasciish.com/> (<https://datasciish.com/>)

1.1 Overview

[...]

Vi(sion) Studios, a new streaming service, is looking to launch a concept called "Digital Cinema Night" where customers can build a "Cinema Night" around specific movies. They've hired us to build a Recommendation System in order to launch this concept.

Data, Methodology, and Analysis: we've explored data from MovieLens which captures Movies, Ratings, Genres, and Year. For this analysis, we used all metrics available with the exception of Tags which may be used for later analysis.

Results & Recommendations: After analyzing data from databases with movies heavily weighted in the 1990s through 2000s, we've built a recommendation system for Vi(sion) Studios' "Digital Cinema Night."

1.2 Business Objective

Build a model that provides top 5 movie recommendations to a user for the best Digital Cinema Night experience.

1.3 Data Overview

MovieLens Dataset

Data explored:

** denotes data used in current analysis

^ denotes data for future analysis

1. MovieLens - Movies**
2. MovieLens - Ratings**
3. MovieLens - Tags^
4. MovieLens - Links (will not be used)

2 Data Exploration, Cleansing, and Preparation

Data Exploration

- Outlined in comments: files explored, how and/or why data was chosen, which files will be used, and which files will be used in future analysis

Data Cleansing

- Checked for duplicates, NaN values; continuously cleansed data as necessary

Data Preparation

- Core variables: Movie Titles, Ratings, Year, Genre
- Merged Movies and Ratings datasets

```
In [1]: # import libraries
# Surprise documentation - Surprise: https://surprise.readthedocs.io/en/sta

import csv
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from collections import defaultdict
from surprise import Dataset
from surprise import Reader
from surprise import accuracy
from surprise.model_selection import train_test_split
from surprise.model_selection import cross_validate
from surprise.model_selection import GridSearchCV
from surprise.model_selection import KFold
from surprise import NormalPredictor
from surprise import NormalPredictor
from surprise import BaselineOnly
from surprise import KNNBasic
from surprise import KNNWithMeans
from surprise import KNNBaseline
from surprise import SVD
from surprise import SVDpp
from surprise import NMF
from surprise import SlopeOne
from surprise import CoClustering

%matplotlib inline
```

executed in 806ms, finished 04:49:14 2022-03-14

▼ 2.1 Data Exploration and Cleansing

In [2]: *# explore movies data*

```
movies = pd.read_csv('movielens/movies.csv')
movies
```

executed in 18ms, finished 04:49:14 2022-03-14

Out[2]:

	movieId		title	genres
0	1		Toy Story (1995)	Adventure Animation Children Comedy Fantasy
1	2		Jumanji (1995)	Adventure Children Fantasy
2	3		Grumpier Old Men (1995)	Comedy Romance
3	4		Waiting to Exhale (1995)	Comedy Drama Romance
4	5		Father of the Bride Part II (1995)	Comedy
...
9737	193581		Black Butler: Book of the Atlantic (2017)	Action Animation Comedy Fantasy
9738	193583		No Game No Life: Zero (2017)	Animation Comedy Fantasy
9739	193585		Flint (2017)	Drama
9740	193587		Bungo Stray Dogs: Dead Apple (2018)	Action Animation
9741	193609		Andrew Dice Clay: Dice Rules (1991)	Comedy

9742 rows × 3 columns

In [3]: *# explore ratings data*

```
ratings = pd.read_csv('movielens/ratings.csv')
ratings.head()
```

executed in 31ms, finished 04:49:14 2022-03-14

Out[3]:

	userId	movieId	rating	timestamp
0	1	1	4.0	964982703
1	1	3	4.0	964981247
2	1	6	4.0	964982224
3	1	47	5.0	964983815
4	1	50	5.0	964982931

```
In [4]: # explore links data - will not be used

links = pd.read_csv('movielens/links.csv')
links.head()
```

executed in 10ms, finished 04:49:14 2022-03-14

Out[4]:

	movied	imdbid	tmdbid
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

```
In [5]: # explore tags data - will not be used for this analysis
```

```
tags = pd.read_csv('movielens/tags.csv')
tags.head()
```

executed in 9ms, finished 04:49:14 2022-03-14

Out[5]:

	userId	movied	tag	timestamp
0	2	60756	funny	1445714994
1	2	60756	Highly quotable	1445714996
2	2	60756	will ferrell	1445714992
3	2	89774	Boxing story	1445715207
4	2	89774	MMA	1445715200

After exploring all datasets, Movies and Ratings datasets - which provides titles, ratings, userid, and genre - will be used

Movies and Ratings datasets will be merged

```
In [6]: # merge movies and ratings data on movieId column
```

```
movies_and_ratings = pd.merge(left=movies,
                               right=ratings,
                               on='movieId')
```

executed in 17ms, finished 04:49:14 2022-03-14

In [7]: *# explore movies and ratings data*

```
movies_and_ratings.info()
```

executed in 16ms, finished 04:49:14 2022-03-14

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100836 entries, 0 to 100835
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   movieId    100836 non-null  int64
 1   title      100836 non-null  object
 2   genres     100836 non-null  object
 3   userId     100836 non-null  int64
 4   rating     100836 non-null  float64
 5   timestamp  100836 non-null  int64
dtypes: float64(1), int64(3), object(2)
memory usage: 5.4+ MB
```

In [8]: *# check for duplicates*

```
movies_and_ratings.duplicated(keep='first').sum()
```

executed in 28ms, finished 04:49:14 2022-03-14

Out[8]: 0

There are zero duplicates!

In [9]: *# check for NaN values*

```
movies_and_ratings.isna().sum()
```

executed in 11ms, finished 04:49:14 2022-03-14

```
Out[9]: movieId      0
        title        0
        genres       0
        userId       0
        rating       0
        timestamp    0
        dtype: int64
```

There are zero NaN values!

In [10]: *# check number of unique MovieIds*

```
movies_and_ratings['movieId'].nunique()
```

executed in 3ms, finished 04:49:14 2022-03-14

Out[10]: 9724

```
In [11]: # check number of unique userIds

movies_and_ratings['userId'].nunique()

executed in 3ms, finished 04:49:14 2022-03-14
```

Out[11]: 610

The dataset contains 9724 unique entries (based on movieId) and 610 unique users. This is a solid start for our dataset.

```
In [12]: # check dataset columns

movies_and_ratings.columns

executed in 3ms, finished 04:49:14 2022-03-14
```

Out[12]: Index(['movieId', 'title', 'genres', 'userId', 'rating', 'timestamp'], dtype='object')

```
In [13]: # add column for year
# parse out year using indexes (-5:-1)
# create a placeholder (1000000) for movies that do not have a year

year = []

for movie in movies_and_ratings['title']:
    year_separated = movie[-5:-1]
    try: year.append(int(year_separated))
    except: year.append(1000000)

movies_and_ratings['year'] = year
movies_and_ratings.head()

executed in 60ms, finished 04:49:14 2022-03-14
```

Out[13]:

	movieId	title	genres	userId	rating	timestamp	year
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	1	4.0	964982703	1995
1	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	5	4.0	847434962	1995
2	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	7	4.5	1106635946	1995
3	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	15	2.5	1510577970	1995
4	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	17	4.5	1305696483	1995

```
In [14]: # determine how many movies did not have a year in the title

print(len(movies_and_ratings[movies_and_ratings['year'] == 1000000]))
```

executed in 7ms, finished 04:49:14 2022-03-14

30

```
In [15]: # eliminate the 30 titles that did not have a year attached

movies_and_ratings = movies_and_ratings.loc[movies_and_ratings['year']
                                             != 1000000]
```

executed in 7ms, finished 04:49:14 2022-03-14

```
In [16]: # clean movie titles using indexes to trim the last 6 characters of each ti

movies_and_ratings['title'] = movies_and_ratings['title'].str[:-7]
movies_and_ratings['title']
```

executed in 31ms, finished 04:49:14 2022-03-14

```
Out[16]: 0          Toy Story
1          Toy Story
2          Toy Story
3          Toy Story
4          Toy Story
...
100831  Black Butler: Book of the Atlantic
100832          No Game No Life: Zero
100833          Flint
100834  Bungo Stray Dogs: Dead Apple
100835  Andrew Dice Clay: Dice Rules
Name: title, Length: 100806, dtype: object
```

```
In [17]: # look at info for cleaned dataset

movies_and_ratings.info()
```

executed in 13ms, finished 04:49:14 2022-03-14

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100806 entries, 0 to 100835
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   movieId     100806 non-null  int64
1   title       100806 non-null  object
2   genres      100806 non-null  object
3   userId      100806 non-null  int64
4   rating      100806 non-null  float64
5   timestamp   100806 non-null  int64
6   year        100806 non-null  int64
dtypes: float64(1), int64(4), object(2)
memory usage: 6.2+ MB
```


In [18]: `movies_and_ratings.head()`

executed in 8ms, finished 04:49:14 2022-03-14

Out[18]:

	movielfld	title	genres	userId	rating	timestamp	year
0	1	Toy Story	Adventure Animation Children Comedy Fantasy	1	4.0	964982703	1995
1	1	Toy Story	Adventure Animation Children Comedy Fantasy	5	4.0	847434962	1995
2	1	Toy Story	Adventure Animation Children Comedy Fantasy	7	4.5	1106635946	1995
3	1	Toy Story	Adventure Animation Children Comedy Fantasy	15	2.5	1510577970	1995
4	1	Toy Story	Adventure Animation Children Comedy Fantasy	17	4.5	1305696483	1995

In [19]: `# drop timestamp column`

```
movies_and_ratings = movies_and_ratings.drop(columns=['timestamp'])
movies_and_ratings
```

executed in 15ms, finished 04:49:14 2022-03-14

Out[19]:

	movielfld	title	genres	userId	rating	year
0	1	Toy Story	Adventure Animation Children Comedy Fantasy	1	4.0	1995
1	1	Toy Story	Adventure Animation Children Comedy Fantasy	5	4.0	1995
2	1	Toy Story	Adventure Animation Children Comedy Fantasy	7	4.5	1995
3	1	Toy Story	Adventure Animation Children Comedy Fantasy	15	2.5	1995
4	1	Toy Story	Adventure Animation Children Comedy Fantasy	17	4.5	1995
...
100831	193581	Black Butler: Book of the Atlantic	Action Animation Comedy Fantasy	184	4.0	2017
100832	193583	No Game No Life: Zero	Animation Comedy Fantasy	184	3.5	2017
100833	193585	Flint	Drama	184	3.5	2017
100834	193587	Bungo Stray Dogs: Dead Apple	Action Animation	184	3.5	2018
100835	193609	Andrew Dice Clay: Dice Rules	Comedy	331	4.0	1991

100806 rows × 6 columns

```
In [20]: # explore years in dataset - groupby year and title

movie_yr_count = movies_and_ratings.groupby('year').count()['title']
                                                    .reset_index()

movie_yr_count.head()
```

executed in 18ms, finished 04:49:14 2022-03-14

Out[20]:

	year	title
0	1902	5
1	1903	2
2	1908	1
3	1915	1
4	1916	5



2.2 Data Visualizations

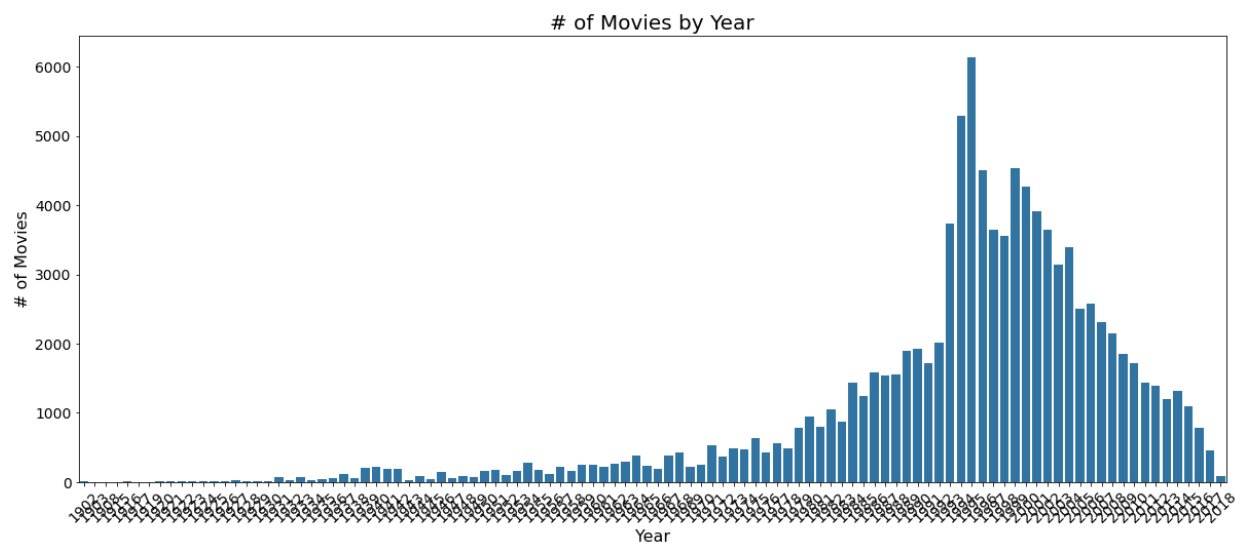
```
In [21]: # plot the years - countplot

fig, ax = plt.subplots(figsize=(20,8))

sns.countplot(x='year',
              data=movies_and_ratings,
              color='tab:blue');

plt.xlabel("Year", fontsize=16)
plt.ylabel("# of Movies", fontsize=16)
plt.title("# of Movies by Year", fontsize=20)
plt.xticks(fontsize=14, rotation=45)
plt.yticks(fontsize=14)
ax.grid(False)
plt.show()
```

executed in 1.22s, finished 04:49:15 2022-03-14

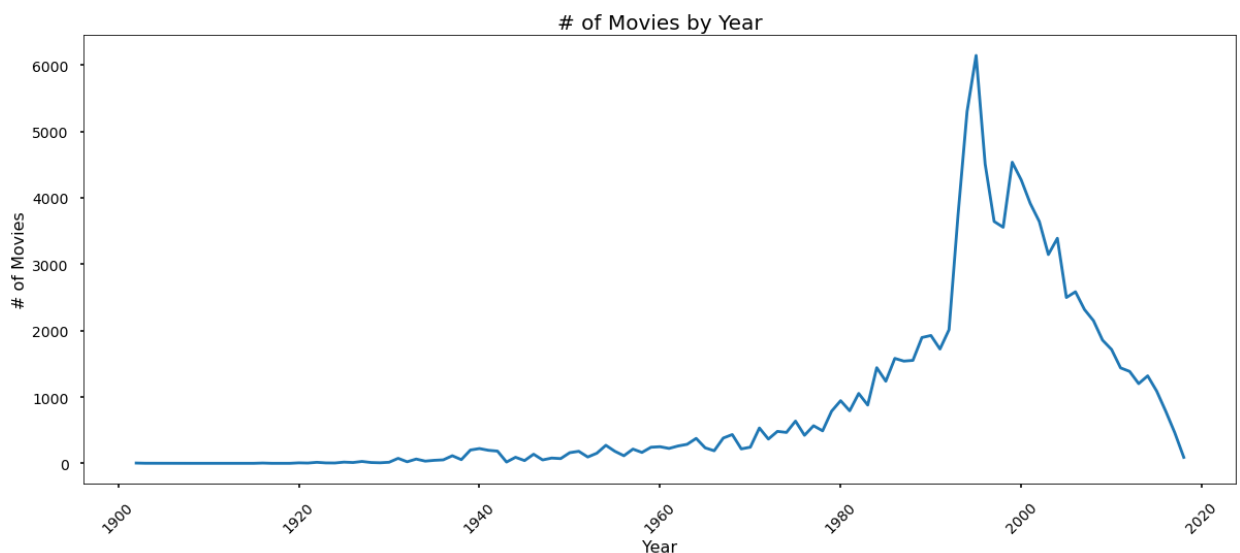


In [22]: *# plot the years - lineplot*

```
with plt.style.context('seaborn-poster'):
    fig, ax = plt.subplots(figsize=(20,8))
    sns.lineplot(x='year',
                  y='title',
                  data=movie_yr_count,
                  color='tab:blue');

    plt.xlabel("Year", fontsize=16)
    plt.ylabel("# of Movies", fontsize=16)
    plt.title("# of Movies by Year", fontsize=20)
    plt.xticks(fontsize=14, rotation=45)
    plt.yticks(fontsize=14)
    ax.grid(False)
    plt.show()
```

executed in 128ms, finished 04:49:15 2022-03-14



In [23]: *# continue exploring the years in movie dataset*
use (normalize=True) so you do not need to do what you did in cell below

```
movies_and_ratings['year'].value_counts(normalize=True)*100
```

executed in 5ms, finished 04:49:15 2022-03-14

```
Out[23]: 1995    6.093883
1994    5.253656
1999    4.498740
1996    4.472948
2000    4.233875
...
1903    0.001984
1908    0.000992
1915    0.000992
1919    0.000992
1917    0.000992
Name: year, Length: 106, dtype: float64
```

```
In [24]: # reference of what not to do
# calculate % of movies by year

movies_years_composition = (
    movies_and_ratings['year'].value_counts() /
    len(movies_and_ratings['year']))*100

movies_years_composition
```

executed in 6ms, finished 04:49:15 2022-03-14

```
Out[24]: 1995      6.093883
1994      5.253656
1999      4.498740
1996      4.472948
2000      4.233875
...
1903      0.001984
1908      0.000992
1915      0.000992
1919      0.000992
1917      0.000992
Name: year, Length: 106, dtype: float64
```

```
In [25]: # look at standard metrics for ratings in dataset
```

```
movies_and_ratings['rating'].describe()
```

executed in 9ms, finished 04:49:15 2022-03-14

```
Out[25]: count      100806.000000
mean           3.501592
std            1.042414
min            0.500000
25%            3.000000
50%            3.500000
75%            4.000000
max            5.000000
Name: rating, dtype: float64
```

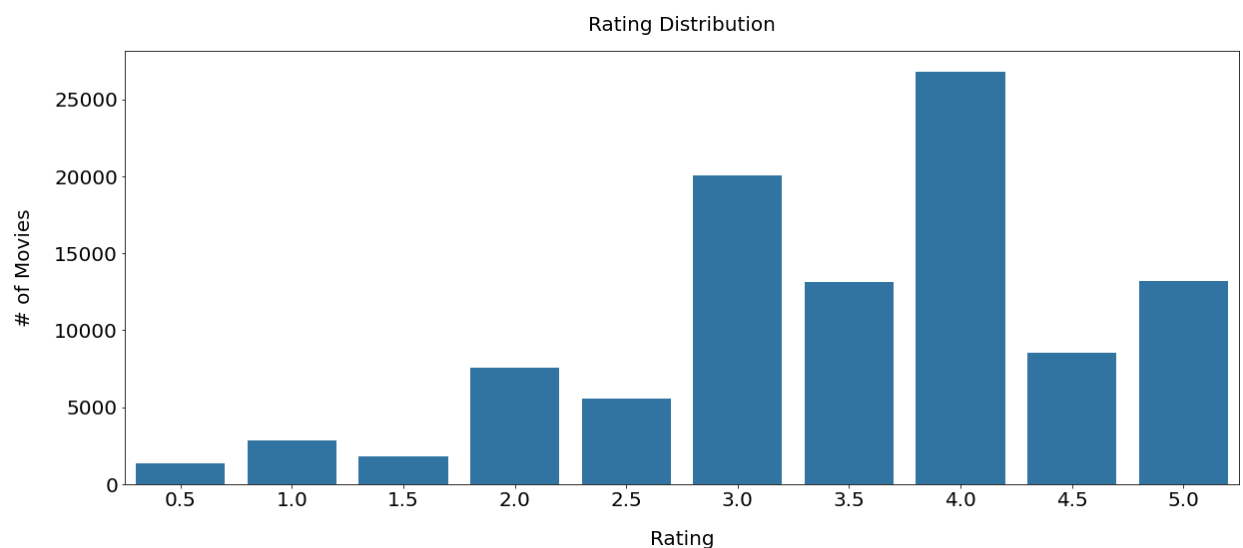
```
In [26]: # plot number of ratings

fig, ax = plt.subplots(figsize=(20,8))

sns.countplot(x='rating',
              data=movies_and_ratings,
              color='tab:blue');

plt.title("Rating Distribution",fontsize=20, pad=20)
plt.xlabel("Rating", fontsize=20, labelpad=20)
plt.ylabel("# of Movies", fontsize=20, labelpad=20)
plt.xticks(fontsize=20)
plt.yticks(fontsize=20)
ax.grid(False)
plt.show()
```

executed in 137ms, finished 04:49:15 2022-03-14



```
In [27]: # CLEAN GENRE DATA
# separate genres from each other and see how many genres there are

genre_df = pd.DataFrame(movies_and_ratings['genres'].
                        str.split('|').
                        tolist(),
                        index=movies_and_ratings['movieId']
                        ).stack()

# reset index

genre_df = genre_df.reset_index([0, 'movieId'])

# create columns movieId and genre

genre_df.columns = ['movieId', 'genre']
```

executed in 243ms, finished 04:49:16 2022-03-14

In [28]: *# explore genre data*

```
genre_df.head()
```

executed in 5ms, finished 04:49:16 2022-03-14

Out[28]:

	movielfld	genre
0	1	Adventure
1	1	Animation
2	1	Children
3	1	Comedy
4	1	Fantasy

In [29]: *# explore genre value counts*

```
genre_df['genre'].value_counts()
```

executed in 39ms, finished 04:49:16 2022-03-14

Out[29]:

Drama	41923
Comedy	39049
Action	30623
Thriller	26446
Adventure	24157
Romance	18124
Sci-Fi	17233
Crime	16679
Fantasy	11831
Children	9207
Mystery	7674
Horror	7287
Animation	6982
War	4858
IMAX	4145
Musical	4138
Western	1930
Documentary	1219
Film-Noir	870
(no genres listed)	38

Name: genre, dtype: int64

In [30]: *# genre data data to explore most popular genres*

```
genre_df_sorted = (  
    genre_df['genre']  
    .value_counts()  
    .sort_values(ascending=False)  
    .reset_index()  
)
```

genre_df_sorted

executed in 33ms, finished 04:49:16 2022-03-14

Out[30]:

	index	genre
0	Drama	41923
1	Comedy	39049
2	Action	30623
3	Thriller	26446
4	Adventure	24157
5	Romance	18124
6	Sci-Fi	17233
7	Crime	16679
8	Fantasy	11831
9	Children	9207
10	Mystery	7674
11	Horror	7287
12	Animation	6982
13	War	4858
14	IMAX	4145
15	Musical	4138
16	Western	1930
17	Documentary	1219
18	Film-Noir	870
19	(no genres listed)	38

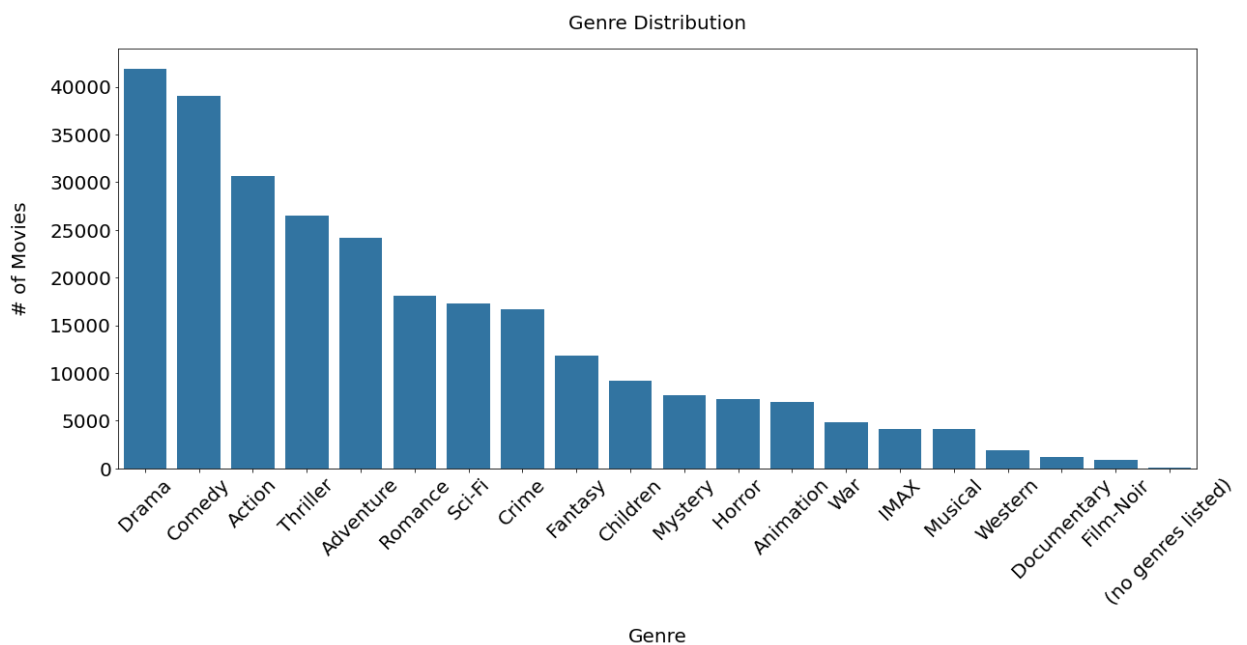
In [31]: *# plot value counts by genre*

```
fig, ax = plt.subplots(figsize=(20,8))

sns.barplot(x='index',
            y='genre',
            data=genre_df_sorted,
            color='tab:blue');

plt.title("Genre Distribution",fontsize=20, pad=20)
plt.xlabel("Genre", fontsize=20, labelpad=20)
plt.ylabel("# of Movies", fontsize=20, labelpad=20)
plt.xticks(fontsize=20, rotation=45)
plt.yticks(fontsize=20)
ax.grid(False)
plt.show()
```

executed in 239ms, finished 04:49:16 2022-03-14



In [32]: *# create movie_popularity variable using value counts*

```
movie_popularity = movies_and_ratings["title"].value_counts()

# create variable "popular_movies" for only movies appearing > 50 times
popular_movies = movie_popularity[movie_popularity > 50]
```

executed in 14ms, finished 04:49:16 2022-03-14

In [33]: *# explore movie_popularity statistics*

```
movie_popularity.describe()
```

executed in 6ms, finished 04:49:16 2022-03-14

Out[33]:

count	9423.000000
mean	10.697867
std	22.837215
min	1.000000
25%	1.000000
50%	3.000000
75%	9.000000
max	329.000000

Name: title, dtype: float64

In [34]: *# explore popular_movies statistics*

```
popular_movies.describe()
```

executed in 5ms, finished 04:49:16 2022-03-14

Out[34]:

count	444.000000
mean	93.189189
std	45.947219
min	51.000000
25%	61.000000
50%	77.000000
75%	108.250000
max	329.000000

Name: title, dtype: float64

In [35]: *# create a column called 'RatingsCounts' - number of ratings for each movie*

```
num_ratings = pd.DataFrame(  
    movies_and_ratings  
    .groupby('movieId')  
    .count()['rating']  
    ).reset_index()  
  
movies_and_ratings = pd.merge(left=movies_and_ratings,  
                               right=num_ratings,  
                               on='movieId')  
  
movies_and_ratings.rename(columns={'rating_x': 'rating',  
                                   'rating_y': 'RatingsCount'},  
                           inplace=True)
```

executed in 65ms, finished 04:49:16 2022-03-14

In [36]: `movies_and_ratings.info()`

executed in 17ms, finished 04:49:16 2022-03-14

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 100806 entries, 0 to 100805
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   movieId         100806 non-null   int64
1   title           100806 non-null   object
2   genres          100806 non-null   object
3   userId          100806 non-null   int64
4   rating          100806 non-null   float64
5   year            100806 non-null   int64
6   RatingsCount    100806 non-null   int64
dtypes: float64(1), int64(4), object(2)
memory usage: 6.2+ MB
```

In [37]: `# sort by ratings counts`

```
movies_and_ratings_sorted = (
    movies_and_ratings
    .sort_values(by='RatingsCount', ascending=False)
    .drop_duplicates('movieId')
)
```

`movies_and_ratings_sorted`

executed in 27ms, finished 04:49:16 2022-03-14

Out[37]:

	movieId	title	genres	userId	rating	year	RatingsCount
10332	356	Forrest Gump	Comedy Drama Romance War	589	5.0	1994	329
8860	318	Shawshank Redemption, The	Crime Drama	400	5.0	1994	317
7886	296	Pulp Fiction	Comedy Crime Drama Thriller	45	5.0	1994	307
16321	593	Silence of the Lambs, The	Crime Horror Thriller	201	5.0	1991	279
45053	2571	Matrix, The	Action Sci-Fi Thriller	80	4.5	1999	278
...
79447	32442	Greedy	Comedy	599	2.5	1994	1
79446	32440	If Looks Could Kill	Action Comedy	599	2.0	1991	1
79445	32392	800 Bullets (800 Balas)	Comedy Crime Drama Western	387	2.5	2002	1
97640	115667	Love, Rosie	Comedy Romance	563	3.5	2014	1
100805	193609	Andrew Dice Clay: Dice Rules	Comedy	331	4.0	1991	1

9701 rows × 7 columns

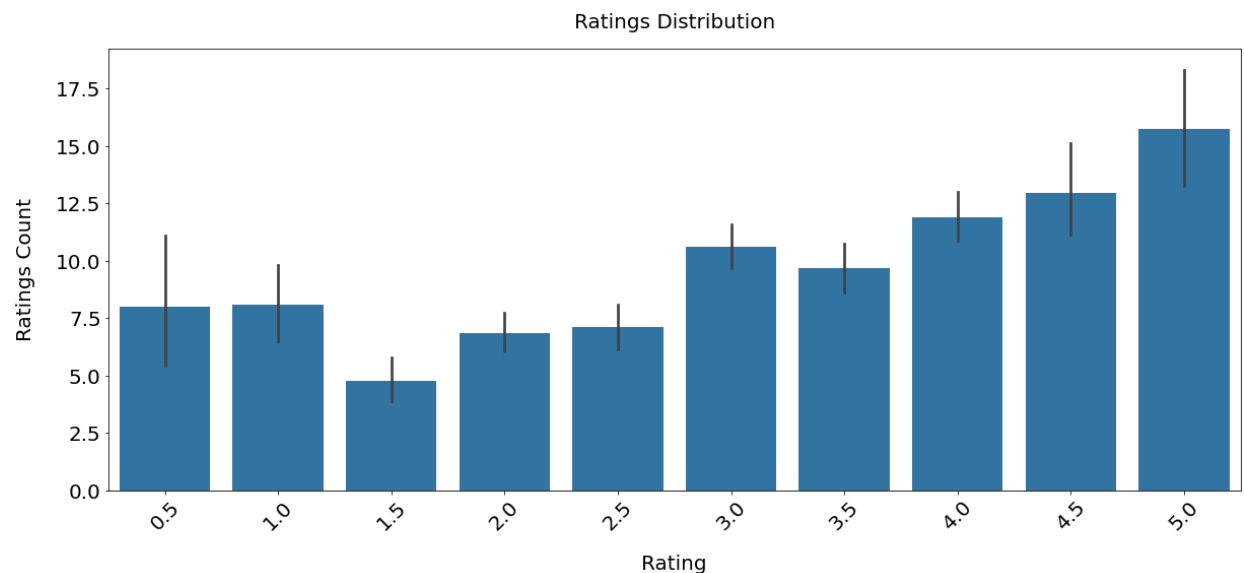
In [38]: *# plot value counts by RatingsCount*

```
fig, ax = plt.subplots(figsize=(20,8))

sns.barplot(x='rating',
            y='RatingsCount',
            data=movies_and_ratings_sorted,
            color='tab:blue');

plt.title("Ratings Distribution",fontsize=20, pad=20)
plt.xlabel("Rating", fontsize=20, labelpad=20)
plt.ylabel("Ratings Count", fontsize=20, labelpad=20)
plt.xticks(fontsize=20, rotation=45)
plt.yticks(fontsize=20)
ax.grid(False)
plt.show()
```

executed in 382ms, finished 04:49:17 2022-03-14



In [39]: *# explore latest dataset*

```
movies_and_ratings_sorted.info()
```

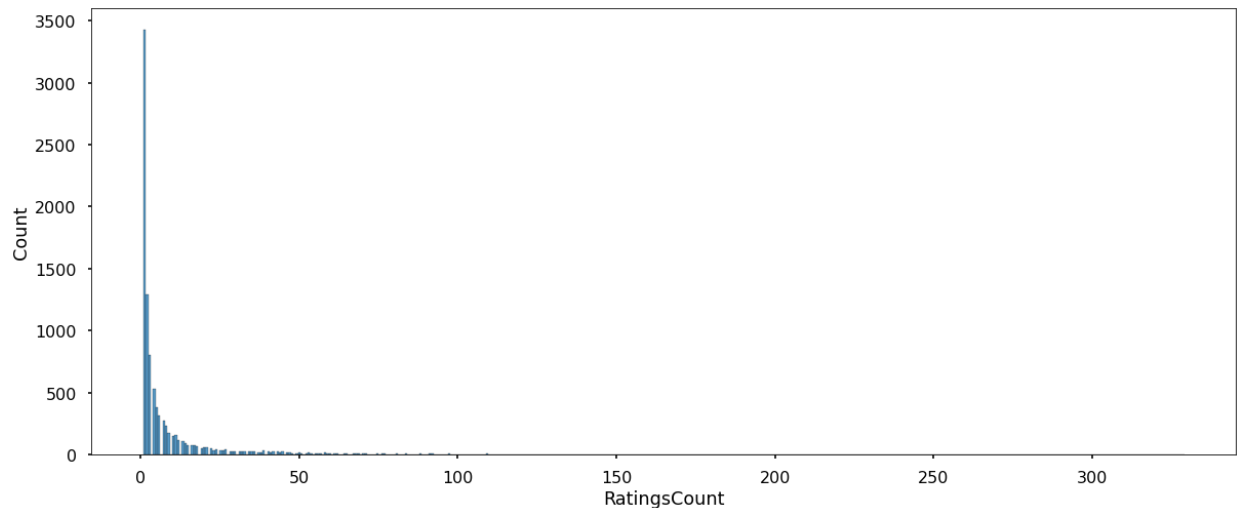
executed in 9ms, finished 04:49:17 2022-03-14

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9701 entries, 10332 to 100805
Data columns (total 7 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   movieId         9701 non-null   int64
1   title           9701 non-null   object
2   genres          9701 non-null   object
3   userId          9701 non-null   int64
4   rating          9701 non-null   float64
5   year            9701 non-null   int64
6   RatingsCount    9701 non-null   int64
dtypes: float64(1), int64(4), object(2)
memory usage: 926.3+ KB
```

In [40]: *# plot ratings counts*

```
with plt.style.context('seaborn-poster'):
    fig, ax = plt.subplots(figsize=(20,8))
    sns.histplot(x='RatingsCount',
                 data=movies_and_ratings_sorted,
                 color='tab:blue');
    ax.grid(False)
```

executed in 563ms, finished 04:49:17 2022-03-14



In [41]: *# explore most rated movies*

```
most Rated movies = movies_and_ratings_sorted[
    movies_and_ratings['RatingsCount']>2]

most Rated movies.info()
```

executed in 9ms, finished 04:49:17 2022-03-14

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4979 entries, 10332 to 53319
Data columns (total 7 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   movieId         4979 non-null   int64
 1   title           4979 non-null   object
 2   genres          4979 non-null   object
 3   userId          4979 non-null   int64
 4   rating          4979 non-null   float64
 5   year            4979 non-null   int64
 6   RatingsCount    4979 non-null   int64
dtypes: float64(1), int64(4), object(2)
memory usage: 311.2+ KB
```

```
<ipython-input-41-91ba2498a61a>:3: UserWarning: Boolean Series key will be
reindexed to match DataFrame index.
```

```
most Rated movies = movies_and_ratings_sorted[
```

3 Build Recommendation System

▼ 3.1 Prepare Data for Modeling

▼ 3.1.1 Read in dataset

```
In [42]: # instantiate a reader

reader = Reader(rating_scale=(0,5))

# load data

movies = Dataset.load_from_df(
    movies_and_ratings_sorted[['userId', 'title', 'rating']],
    reader)
```

executed in 11ms, finished 04:49:17 2022-03-14

▼ 3.1.2 Train Test Split

```
In [43]: # create train and test datasets

train, test = train_test_split(movies, test_size=.2)
```

executed in 13ms, finished 04:49:17 2022-03-14

```
In [44]: # look at train data
```

```
train
```

executed in 3ms, finished 04:49:17 2022-03-14

```
Out[44]: <surprise.trainset.Trainset at 0x7fde7973e940>
```


In [47]: *# run 5-fold cross_validation*

```
cross_validate(svd, movies, measures=['rmse', 'mae'], cv=5, verbose=True)
```

executed in 2.08s, finished 04:49:20 2022-03-14

Evaluating RMSE, MAE of algorithm SVD on 5 split(s).

	Fold 1	Fold 2	Fold 3	Fold 4	Fold 5	Mean	Std
RMSE (testset)	0.9595	0.9374	0.9349	0.9351	0.9409	0.9415	0.009
2							
MAE (testset)	0.7455	0.7323	0.7266	0.7288	0.7368	0.7340	0.006
7							
Fit time	0.37	0.41	0.38	0.38	0.37	0.38	0.01
Test time	0.01	0.01	0.01	0.01	0.06	0.02	0.02

Out[47]: {'test_rmse': array([0.95946107, 0.93736966, 0.93491585, 0.93505028, 0.94094627]),
 'test_mae': array([0.74553923, 0.73226289, 0.72662335, 0.72877489, 0.73680258]),
 'fit_time': (0.37498903274536133,
 0.4071052074432373,
 0.38315701484680176,
 0.3824949264526367,
 0.37239694595336914),
 'test_time': (0.0072138309478759766,
 0.007066965103149414,
 0.007613182067871094,
 0.010246038436889648,
 0.0639491081237793)}

In [48]: *# GRIDSEARCH - SVD Model*

```
param_grid = {'n_epochs': [5, 10], 'lr_all': [0.002, 0.005],  

  'reg_all': [0.4, 0.6]}
```

```
gs_svd = GridSearchCV(SVD,  

  param_grid,  

  measures=['rmse', 'mae'],  

  cv=5)
```

```
gs_svd.fit(movies)
```

best RMSE score

```
print(gs_svd.best_score['rmse'])
```

combination of parameters that gave the best RMSE score

```
print(gs_svd.best_params['rmse'])
```

executed in 6.27s, finished 04:49:26 2022-03-14

```
0.9579050537876668
```

```
{'n_epochs': 10, 'lr_all': 0.005, 'reg_all': 0.4}
```


In [49]: *# FINAL SVD MODEL*

```
svd_algo = gs_svd.best_estimator['rmse']
svd_algo.fit(movies.build_full_trainset())
```

```
svd_best_predictions = svd_algo.test(test)
```

check RMSE and MAE of final SVD model

```
accuracy.rmse(svd_best_predictions)
accuracy.mae(svd_best_predictions)
```

executed in 262ms, finished 04:49:26 2022-03-14

RMSE: 0.8601

MAE: 0.6738

Out[49]: 0.6737755408300745

Final SVD Model Performance: average predictions are 0.86 stars away from the actual rating

View five predictions

r_ui: actual rating

est: estimated rating

In [50]: *# check five predictions*

```
svd_best_predictions[:5]
```

executed in 4ms, finished 04:49:26 2022-03-14

Out[50]: [Prediction(uid=474, iid='Wildcats', r_ui=1.0, est=3.2312788476843943, details={'was_impossible': False}),
 Prediction(uid=534, iid='21 Jump Street', r_ui=4.0, est=3.636737034488619, details={'was_impossible': False}),
 Prediction(uid=567, iid="Ivan's Childhood (a.k.a. My Name is Ivan) (Ivan ovo detstvo)", r_ui=0.5, est=2.2375535472515216, details={'was_impossible': False}),
 Prediction(uid=113, iid='Saving Grace', r_ui=5.0, est=3.692164164558661, details={'was_impossible': False}),
 Prediction(uid=91, iid='Enter the Dragon', r_ui=4.0, est=3.34316976234394, details={'was_impossible': False})]

Test/make a prediction with user and movie id

```
In [51]: # get prediction for specific user and item

user_id = str(111)
movie_id = str(111)

svd_predictor = svd_algo.predict(user_id, movie_id, r_ui=4, verbose=True)

executed in 3ms, finished 04:49:26 2022-03-14

user: 111      item: 111      r_ui = 4.00    est = 3.27    {'was_imposs
ible': False}
```

For reference: function provided by Surprise to find top_n recommendations for users

To be improved in Final Concept Launch

```
In [52]: from collections import defaultdict

def get_top_n(svd_best_predictions, n=10):
    """Return the top-N recommendation for each user from a set of predictions

    Args:
        predictions(list of Prediction objects): The list of predictions, a
            returned by the test method of an algorithm.
        n(int): The number of recommendation to output for each user. Default
            is 10.

    Returns:
        A dict where keys are user (raw) ids and values are lists of tuples:
            [(raw item id, rating estimation), ...] of size n.
    """

    # Map the predictions to each user
    top_n = defaultdict(list)
    for uid, iid, true_r, est, _ in svd_best_predictions:
        top_n[uid].append((iid, est))

    # Sort the predictions for each user and retrieve the k highest ones
    for uid, user_ratings in top_n.items():
        user_ratings.sort(key=lambda x: x[1], reverse=True)
        top_n[uid] = user_ratings[:n]

    return top_n

top_n = get_top_n(svd_best_predictions, n=10)

# Print the recommended items for each user
for uid, user_ratings in top_n.items():
    print(uid, [iid for (iid, _) in user_ratings])
```

executed in 30ms, finished 04:49:26 2022-03-14

```
474 ['Mildred Pierce', 'Safety Last!', 'Auntie Mame', 'Jane Eyre', 'Ste
amboat Bill, Jr.', 'Anna Karenina', 'Divided We Fall (Musíme si pomáha
t)', 'Dark Water (Honogurai mizu no soko kara)', 'Zelary', 'Dark Passag
e']
534 ['10th Kingdom, The', '21 Jump Street', 'Oz the Great and Powerfu
l', 'Doomsday', 'Absolutely Anything', 'Lone Ranger, The', 'Sinbad: Leg
end of the Seven Seas', '47 Ronin', 'Stalingrad', 'Star Wars: Episode I
I - Attack of the Clones']
567 ['Eye in the Sky', 'Frances Ha', 'Everest', 'Only Yesterday (Omohid
e poro poro)', 'Lilya 4-Ever (Lilja 4-ever)', 'The Martian', 'Fences',
"The Devil's Candy", 'Jack and Jill', "God's Not Dead"]
113 ["For Roseanna (Roseanna's Grave)", 'Crying Game, The', 'Saving Gra
ce', 'Morning After, The', 'Living Out Loud']
91 ['Pirates of Silicon Valley', 'Enter the Dragon', 'Body Snatcher, Th
e', 'Martin', 'Robot Carnival (Roboto kânibauru)', 'Startup.com', 'Craz
ies, The (a.k.a. Code Name: Trixie)', 'Fist of Fury (Chinese Connectio
n, The) (Jing wu men)', 'Outlaw Josey Wales, The', 'Godzilla, King of t
he Monsters! (Kaijû-ô Gojira)']
448 ['42', 'City Island', 'Florence Foster Jenkins', 'Alan Partridge: A
...']
```



3.2.2 KNN Baseline Model (k-Nearest Neighbor)

k-Nearest Neighbor Baseline Model

Collaborative filtering algorithm that uses Baseline rating

```
In [53]: # instantiate KNN Baseline model
knnb = KNNBaseline(k=50)

# fit data on train set
knnb.fit(train)

# predict/test on test aset
knnb_predictions = knnb.test(test)

# check RMSE and MAE
accuracy.rmse(knnb_predictions)
accuracy.mae(knnb_predictions)
```

executed in 43ms, finished 04:49:26 2022-03-14

```
Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
RMSE: 0.9607
MAE: 0.7492
```

Out[53]: 0.7491622763262366

Starting KNN Baseline Model Performance: average predictions are 0.96 stars away from the actual rating

In [54]: # GRIDSEARCH - KNN Baseline Model

```
param_grid = {'k': [10, 20, 30, 40, 50, 60, 70, 80, 90, 100],
              'sim_options': {'user_based': [True, False]},\
              'bsl_options': {'method': ['als', 'sgd']}}

gs_knnb = GridSearchCV(KNNBaseline,
                       param_grid,
                       measures=['rmse', 'mae'],
                       cv=5)

gs_knnb.fit(movies)

print(gs_knnb.best_score['rmse'])
print(gs_knnb.best_params['rmse'])
```

```
executed in 1m 42.6s, finished 04:51:09 2022-03-14
Estimating biases using sgd...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using sgd...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using sgd...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using als...
Computing the msd similarity matrix...
Done computing similarity matrix.
Estimating biases using als...
Computing the msd similarity matrix...
```

In [55]: *# FINAL KNN MODEL*

```
knnb_algo = gs_knnb.best_estimator['rmse']
knnb_algo.fit(movies.build_full_trainset())

knnb_best_predictions = knnb_algo.test(test)

# check RMSE and MAE of final KNN model

accuracy.rmse(knnb_best_predictions)
accuracy.mae(knnb_best_predictions)
```

executed in 2.39s, finished 04:51:11 2022-03-14

```
Estimating biases using SGD...
Computing the msd similarity matrix...
Done computing similarity matrix.
RMSE: 0.3926
MAE: 0.2647
```

Out[55]: 0.2646965200335826

Final KNN Baseline Model Performance: average predictions are 0.39 stars away from the actual rating

View five predictions

r_ui: actual rating
est: estimated rating

In [56]: knnb_best_predictions[:5]

executed in 3ms, finished 04:51:11 2022-03-14

Out[56]: [Prediction(uid=474, iid='Wildcats', r_ui=1.0, est=1.8593421262239793, details={'actual_k': 100, 'was_impossible': False}),
Prediction(uid=534, iid='21 Jump Street', r_ui=4.0, est=3.8572007627783793, details={'actual_k': 44, 'was_impossible': False}),
Prediction(uid=567, iid='Ivan's Childhood (a.k.a. My Name is Ivan) (Ivan ovo detstvo)', r_ui=0.5, est=1.007962576668691, details={'actual_k': 88, 'was_impossible': False}),
Prediction(uid=113, iid='Saving Grace', r_ui=5.0, est=4.6121700668745635, details={'actual_k': 21, 'was_impossible': False}),
Prediction(uid=91, iid='Enter the Dragon', r_ui=4.0, est=3.7147285147535354, details={'actual_k': 67, 'was_impossible': False})]



3.2.3 NMF Model (Non-Negative Matrix Factorization)

Non-Negative Matrix Factorization Model

Collaborative filtering model based on Non-Negative Matrix Factorization

```
In [57]: # instantiate NMF model
nmf = NMF()

# fit data on train set
nmf.fit(train)

# predict/test on test aset
nmf_predictions = nmf.test(test)

# check RMSE and MAE

accuracy.rmse(nmf_predictions)
accuracy.mae(nmf_predictions)
```

executed in 831ms, finished 04:51:12 2022-03-14

RMSE: 1.1151
MAE: 0.8960

Out[57]: 0.8960196508519168

Starting NMF Model Performance: average predictions are 1.12 stars away from the actual rating

```
In [58]: # GRIDSEARCH - NMF Model

param_grid = {'n_factors': [1,2,3,4,5,6,7,8,9,10],
              'n_epochs': [100],
              'biased': [True],
              'reg_bu': [0.1],
              'reg_bi': [0.1]}

gs_nmf = GridSearchCV(NMF, param_grid, measures=['rmse', 'mae'], cv=5)

gs_nmf.fit(movies)

#nmfb = gs_nmf.best_estimator['rmse']
print(gs_nmf.best_score['rmse'])
print(gs_nmf.best_params['rmse'])
```

executed in 44.4s, finished 04:51:57 2022-03-14

1.0967863468536287
{'n_factors': 1, 'n_epochs': 100, 'biased': True, 'reg_bu': 0.1, 'reg_bi': 0.1}

In [59]: *# FINAL NMF MODEL*

```
nmf_algo = gs_nmf.best_estimator['rmse']
nmf_algo.fit(movies.build_full_trainset())
```

```
nmf_best_predictions = nmf_algo.test(test)
```

check RMSE and MAE of final NMF model

```
accuracy.rmse(nmf_best_predictions)
accuracy.mae(nmf_best_predictions)
```

executed in 817ms, finished 04:51:57 2022-03-14

RMSE: 0.3064

MAE: 0.1708

Out[59]: 0.1707548649981007

Final NMF Model Performance: average predictions are 0.30 stars away from the actual rating

View five predictions

r_ui: actual rating

est: estimated rating

In [60]: nmf_best_predictions[:5]

executed in 3ms, finished 04:51:57 2022-03-14

Out[60]: [Prediction(uid=474, iid='Wildcats', r_ui=1.0, est=2.216686139297967, details={'was_impossible': False}),
 Prediction(uid=534, iid='21 Jump Street', r_ui=4.0, est=3.9307526989313994, details={'was_impossible': False}),
 Prediction(uid=567, iid="Ivan's Childhood (a.k.a. My Name is Ivan) (Ivan ovo detstvo)", r_ui=0.5, est=1.1027548082478766, details={'was_impossible': False}),
 Prediction(uid=113, iid='Saving Grace', r_ui=5.0, est=4.895730843115939, details={'was_impossible': False}),
 Prediction(uid=91, iid='Enter the Dragon', r_ui=4.0, est=3.9308589487980776, details={'was_impossible': False})]

4 Evaluation and Conclusions

We've built a Recommendation System to execute on Vi(sion) Studios "Digital Cinema Night" concept!

Model Evaluation

- Singular Vector Decomposition Model (SVD):
 - * Pre-Tuned RMSE: 0.9502
 - * Tuned RMSE: 0.8622

- k-Nearest Neighbor Baseline Model (KNNB):
 - * Pre-Tuned RMSE: 0.9604
 - * Tuned RMSE: 0.3863
- Non-Negative Matrix Factorization Model (NMF):
 - * Pre-Tuned RMSE: 1.1232
 - * Tuned RMSE: 0.2964

We've decided to use the SVD Model to launch "Digital Cinema Night." While all models showed improvement after tuning and KNN Baseline and NMF Models have low RMSEs (predictions are close to actual ratings), in order to reduce the risk of over-fitting, we will use the SVD model. With an RMSE of 0.86, we feel best our recommendation system make more expansive recommendations to users.

Summary of recommendations

- Use SVD Model to launch
- Test for 6 months and re-evaluate
 - *Revisit KNN Baseline and/or NMF Models if necessary
- Launch concept to target Millennial audience
 - *Movies weighted toward movies from 1990s to early 2000s
 - *Build early loyalty with influential demographic

Further considerations:

- Explore launching next round with genre and year selections as a function

5 Future Work

This is just the beginning!

Future work:

- Integrate genre and year functionality into recommendation system
- Launch website allowing users to build "Digital Cinema Night" by typing in movie names, genre, and/or year

5.1 APPENDIX

Output to explore and aspire toward

In [61]: *# cross-validation results dataframe*

```
results_df = pd.DataFrame.from_dict(gs_svd.cv_results)  
results_df
```

executed in 97ms, finished 04:51:58 2022-03-14

Out[61]:

	split0_test_rmse	split1_test_rmse	split2_test_rmse	split3_test_rmse	split4_test_rmse	mean_test
0	0.977595	0.989719	0.979794	0.982992	0.995029	0.9
1	0.984173	0.995288	0.984269	0.989404	0.999870	0.9
2	0.961932	0.971756	0.961698	0.964057	0.975037	0.9

In [62]: *# create matrix for movies and ratings*

```
matrix = movies_and_ratings.pivot_table(
    index='userId',
    columns='title',
    values='rating'
)

matrix.head()
```

executed in 172ms, finished 04:51:58 2022-03-14

Out[62]:

	title	'71	'Hellboy': The Seeds of Creation	'Round Midnight	'Salem's Lot	'Til There Was You	'Tis the Season for Love	'burbs, The	'night Mother	(500) Days of Summer	*batteries not included	...
userId												
1	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	..
2	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	..
3	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	..
4	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	..
5	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	..

5 rows × 9423 columns

In [63]: *# create recommendation system to make 5 recommendations to users*

```
def cinema_night_for(title, ratings_count_filter=100, number_recommendation
    similar = matrix.corrwith(matrix[title])
    corr_similar = pd.DataFrame(similar, columns=['correlation'])
    corr_similar.dropna(inplace=True)

    orig = movies_and_ratings.copy()

    corr_with_movie = pd.merge(
        left=corr_similar,
        right=orig,
        on='title')[['title', 'correlation', 'RatingsCount']].drop_duplicates()

    result = corr_with_movie[corr_with_movie['RatingsCount'] > ratings_count_filter]

    return result.head(number_recommendations)
```

executed in 4ms, finished 04:51:58 2022-03-14

In [64]: `cinema_night_for('Toy Story')`

executed in 1.75s, finished 04:51:59 2022-03-14

```
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2526: RuntimeWarning: Degrees of freedom <= 0 for slice
  c = cov(x, y, rowvar)
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2455: RuntimeWarning: divide by zero encountered in true_divide
  c *= np.true_divide(1, fact)
```

Out[64]:

	title	correlation	RatingsCount
4485	Toy Story	1.000000	215
2174	Incredibles, The	0.643301	125
1505	Finding Nemo	0.618701	141
138	Aladdin	0.611892	183
2914	Monsters, Inc.	0.490231	132
2960	Mrs. Doubtfire	0.446261	144

In [65]: `# TEST RECOMMENDATION SYSTEM using Toy Story`

`cinema_night_for('Toy Story')`

executed in 1.79s, finished 04:52:01 2022-03-14

```
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2526: RuntimeWarning: Degrees of freedom <= 0 for slice
  c = cov(x, y, rowvar)
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2455: RuntimeWarning: divide by zero encountered in true_divide
  c *= np.true_divide(1, fact)
```

Out[65]:

	title	correlation	RatingsCount
4485	Toy Story	1.000000	215
2174	Incredibles, The	0.643301	125
1505	Finding Nemo	0.618701	141
138	Aladdin	0.611892	183
2914	Monsters, Inc.	0.490231	132
2960	Mrs. Doubtfire	0.446261	144

In [66]: *# Test for The Shawshank Redemption*

```
cinema_night_for('Shawshank Redemption, The')
```

executed in 1.76s, finished 04:52:03 2022-03-14

```
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2526: RuntimeWarning: Degrees of freedom <= 0 for slice
  c = cov(x, y, rowvar)
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2455: RuntimeWarning: divide by zero encountered in true_divide
  c *= np.true_divide(1, fact)
```

Out[66]:

	title	correlation	RatingsCount
3890	Shawshank Redemption, The	1.000000	317
1598	Four Weddings and a Funeral	0.446212	103
3787	Schindler's List	0.402202	220
4616	Usual Suspects, The	0.394294	204
3175	Ocean's Eleven	0.391546	119
1838	Green Mile, The	0.382818	111

In [67]: `cinema_night_for('Godfather, The')`

executed in 1.75s, finished 04:52:05 2022-03-14

```
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2526: RuntimeWarning: Degrees of freedom <= 0 for slice
  c = cov(x, y, rowvar)
/Users/v/opt/anaconda3/envs/learn-env/lib/python3.8/site-packages/numpy/lib/function_base.py:2455: RuntimeWarning: divide by zero encountered in true_divide
  c *= np.true_divide(1, fact)
```

Out[67]:

	title	correlation	RatingsCount
1720	Godfather, The	1.000000	192
1721	Godfather: Part II, The	0.782643	129
3740	Schindler's List	0.456661	220
1475	Fight Club	0.445205	218
3718	Saving Private Ryan	0.441377	188
1755	Goodfellas	0.439937	126

```
In [68]: from itertools import permutations

# Create the function to find all permutations
def find_movie_pairs(x):
    pairs = pd.DataFrame(list(permutations(x.values, 2)),
                          columns=['movie_a', 'movie_b'])

    return pairs

# Apply the function to the title column and reset the index
movie_combinations = movies_and_ratings.groupby('userId')['title'].apply(
    find_movie_pairs).reset_index(drop=True)

print(movie_combinations)
```

executed in 17.4s, finished 04:52:22 2022-03-14

	movie_a	movie_b
0	Toy Story	Grumpier Old Men
1	Toy Story	Heat
2	Toy Story	Seven (a.k.a. Se7en)
3	Toy Story	Usual Suspects, The
4	Toy Story	From Dusk Till Dawn
...
60731789	The Fate of the Furious	Rogue One: A Star Wars Story
60731790	The Fate of the Furious	Split
60731791	The Fate of the Furious	John Wick: Chapter Two
60731792	The Fate of the Furious	Get Out
60731793	The Fate of the Furious	Logan

[60731794 rows x 2 columns]

```
In [69]: # calculate how often each item in movie_a occurs with the items in movie_b
combination_counts = movie_combinations.groupby(['movie_a', 'movie_b']).size

# convert the results to a DataFrame and reset the index
combination_counts_df = combination_counts.to_frame(name='size').reset_index()
print(combination_counts_df.head())
```

executed in 18.6s, finished 04:52:41 2022-03-14

	movie_a	movie_b	size
0	'71	(500) Days of Summer	1
1	'71	10 Cloverfield Lane	1
2	'71	127 Hours	1
3	'71	13 Assassins (Jûsan-nin no shikaku)	1
4	'71	13 Hours	1

```
In [70]: # calculate ratings by genre

values = defaultdict(list)
for ind, row in movies_and_ratings.iterrows():
    for genre in row['genres'].split('|'):
        values[genre].append(row['rating'])

genre_list, rating_list = [], []
for key, item in values.items():
    if key not in [0, 1]:
        genre_list.append(key)
        rating_list.append(np.mean(item))

genres_and_ratings = pd.DataFrame([genre_list, rating_list]).T
genres_and_ratings.columns = ['genre', 'avg_rating']
```

executed in 6.39s, finished 04:52:47 2022-03-14

In [71]: genres_and_ratings

executed in 6ms, finished 04:52:47 2022-03-14

Out[71]:

	genre	avg_rating
0	Adventure	3.50886
1	Animation	3.62998
2	Children	3.41311
3	Comedy	3.38471
4	Fantasy	3.491
5	Romance	3.50651
6	Drama	3.65619
7	Action	3.44798
8	Crime	3.65849
9	Thriller	3.49376
10	Horror	3.25834
11	Mystery	3.63246
12	Sci-Fi	3.45601
13	War	3.80815
14	Musical	3.56368
15	Documentary	3.79779
16	IMAX	3.61834
17	Western	3.58394
18	Film-Noir	3.92011
19	(no genres listed)	3.42105